# df++: A Secure JVM Language and Runtime with OOP, Functional, and Declarative Ideas

(Design Review & Implementation Proposal)

Independent Research
contact@example.org

Abstract—This document is a proposal. df++ is a small, statically typed language that combines (i) object-oriented types via interfaces and composition (no inheritance), (ii) functional iteration over List<T>, Set<T>, and Map<K,V> with fusible combinators and first-class pattern matching, and (iii) declarative tasks with contracts (pre/pos) and orchestration (run('...').then('...'), parallel{'...', '...'}. Programs compile to JVM bytecode (JDK 21) via ANTLR+ASM and always execute inside dfVM, a mandatory userspace sandbox that gates I/O by capability, supports deterministic replay, and emits syscall traces. v1 includes a lightweight static contract solver (interval + cardinality + SAT) and basic local type inference. Tooling (CLI/LSP/test/packaging) and ecosystem elements are also proposals in this document.

Index Terms—static typing, pattern matching, local type inference, multiple interfaces, functional collections, declarative tasks, contracts, capability sandbox, deterministic replay, JVM bytecode

#### I. MOTIVATION (PROPOSAL)

One language, three styles. Many systems need OOP modularity, functional clarity, and declarative orchestration. df++ blends these with a deliberately small surface and predictable semantics.

### Design goals.

- **Readable**: interface-based OOP (composition, no inheritance), uniform collections, and explicit tasks.
- **Safe by default**: capability-gated I/O, runtime contracts, 6 reproducible runs.
- Portable: JVM bytecode, Java interop that feels native.
- Auditable: syscall traces + contracts explain failures and performance.
- Implementable now: ANTLR+ASM compiler and a practical sandbox (dfVM) without kernel tricks.

Why dfVM is mandatory. The sandbox delivers least privilege, determinism (virtual time/RNG; HTTP fixtures), resource quotas, and an effect ledger. It replaces many container use-cases with faster startup and lower operational weight while keeping JVM portability.

#### II. LANGUAGE MODEL (V1, PROPOSAL)

### A. Surface

Eight keywords: module, import, type, fn, task, const, let, mut. Identifiers can be normal or back-ticked (e.g., 'foo bar'); task names use single-quoted strings (e.g., 'fetch posts'). Expression-oriented with ternary c ? a : b. No surface for/while loops; use functional iteration.

### B. Types and Interfaces

**Records**: { field = expr, ... }. **ADTs**: type Option<T> = None | Some(T), type Result<T,E> = Ok(T) | Err(E) and user-defined sums. **Generics**: first-order (List<T>, Set<T>, Map<K,V>, Result<T,E>). **Interfaces**: named method sets; values can implement multiple interfaces. Favor composition; no inheritance.

# C. Collections & Pipelines

List<T>: map, flatMap, filter, fold, reduce, scan, zip, take, drop, groupBy, sortBy, forEach, size, isEmpty, nonEmpty; List.range(start,endExcl). Set<T>: union, intersect, diff, contains, size, isEmpty, map, filter, toList. Map<K,V>: get, put, remove, containsKey, keys, values, entries, mapValues, filterKeys, size, isEmpty. Pipelines over these are fusible where semantics allow.

# D. Pattern Matching (v1)

A match-expression is first-class in v1:

```
::df++
type Option<T> = None | Some(T)

fn greet(opt: Option<String>): String =
  match (opt) {
    Some(s) -> "Hello, " + s
    None -> "Hello"
}
```

Patterns in v1 include ADT constructors with positional binders, literals (true/false/null, numbers, strings), record key subsets ({k=v,...} with fresh binders), and wildcard \_. Matching is exhaustive-checked when the domain is finite and known (e.g., Option, small enums); otherwise a warning is issued and a runtime check inserted.

### E. Tasks and Orchestration

A task is task 'name' { pre P; act {  $\dots$  }; pos Q } with pure Boolean pre/pos and effectful act. Orchestration uses:

- run('A').then('B') to sequence.
- parallel { 'A', 'B', 'C' } to fan out on virtual threads and join.

**Testing keyword (tooling proposal)**: test 'name' { arrange; act; assert } recognized by tools; ignored by codegen.

#### F. Contracts, Invariants, and Static Solver (v1)

Contracts run at runtime and feed the trace. A static pass tries to discharge them ahead of time.

#### Contract language. Propositional formulas over:

- numeric constraints: a <b, a  $\leq b, linear forms$
- collection cardinalities: size(xs), isEmpty(xs), subset relations where statically known,
- equality/disequality over literals and finite ADTs.

#### Static solver.

- **Front-end**: normalize to negation-normal form; constant fold; extract known sizes from literals and closed operators (map, filter, take, drop, union, intersect, diff).
- Intervals: compute interval abstractions for Int/Float; propagate through linear ops; detect contradictions.
- Cardinality: track |List|, |Set|, |Map| with flow-sensitive bounds.
- SAT core: CNF-encode residual propositional structure; attempt DPLL; learn simple clauses from interval/cardinality conflicts.
- **Results**: *Proved* (runtime check elided), *Refuted* (counterex-1 ample provided, optional compile error), *Unknown* (runtime 2 check retained).

*Roadmap*: SMT backend (Z3) for quantified finite checks; LTL monitor synthesis for small temporal specs.

# G. Type Inference

#### Local inference:

- let x = expr infers x from expr.
- Lambda parameters may omit types when context suffices.
- Top-level fn signatures recommended (required for exports).

# H. Java Interop (Simplified, Proposal)

Direct, typed imports of JVM classes; usage looks like normal method calls:

# Binding.

- The compiler queries classfiles on the classpath, maps
   JVM types to df++ (int->Int, byte[]->Bytes,
   String->String, generics erased with optional annotations), and generates call sites.
- dfVM validates manifests: each imported class/method must appear in dfpkg.json under jvmAllowlist.

 Reflection is denied; linkage is static; conversions are explicit and well-defined.

III. COMPILER AND RUNTIME (PROPOSAL)

# A. ANTLR $\rightarrow$ ASM Pipeline

**Parsing.** ANTLR4 builds an AST with source spans (Java 21 records).

**Typing.** Nominal types (records/ADTs), first-order generics, interface conformance, exhaustiveness hints for match, local inference, flow non-null.

**Lowering.** Expression IR; closure conversion; pipeline fusion; tail-call to loops; match desugars to decision DAG with guards and binds.

#### Bytecode emission (ASM).

- One public module class: pkg/Module, with static fields for globals.
- Each fn → public static Object f(Object[]) (+ a typed wrapper when exporting to Java).
- Each task t: t\$pre():Z, t\$act():Object, t\$pos():Z.
- COMPUTE\_FRAMES and line/local tables for debugging.

### Walkthrough (example).

```
::df++
// Source
fn add(a: Int, b: Int): Int = a + b
```

#### Lowering sketch:

- 1) AST: Fn(name=add, params=[a:Int,b:Int], ret=Int,
  body=Plus(Var a, Var b))
- 2) IR: Load 0; Load 1; IAdd
- 3) ASM:

```
/* Pseudo-ASM (Java) */
mv = cw.visitMethod(ACC_PUBLIC|ACC_STATIC, "add$0",
  "([Ljava/lang/Object;)Ljava/lang/Object;", null,
    null);
mv.visitCode();
mv.visitVarInsn(ALOAD, 0);
                                       // args
mv.visitInsn(ICONST_0);
mv.visitInsn(AALOAD);
                                       // args[0]
mv.visitTypeInsn(CHECKCAST, "java/lang/Integer");
mv.visitMethodInsn(INVOKEVIRTUAL, "java/lang/Integer",
    intValue", "()I", false);
mv.visitVarInsn(ALOAD, 0);
                                       // args
mv.visitInsn(ICONST_1);
mv.visitInsn(AALOAD);
                                       // args[1]
mv.visitTypeInsn(CHECKCAST, "java/lang/Integer");
mv.visitMethodInsn(INVOKEVIRTUAL, "java/lang/Integer",
    intValue", "()I", false);
mv.visitInsn(IADD);
mv.visitMethodInsn(INVOKESTATIC, "java/lang/Integer",
    valueOf", "(I)Ljava/lang/Integer;", false);
mv.visitInsn(ARETURN);
mv.visitMaxs(0,0);
                                       // COMPUTE_FRAMES
mv.visitEnd();
```

#### **Match lowering.** For

18

```
match (opt) { Some(x) -> f(x); None -> g() }
```

the compiler emits a typed tag-switch (synthetic ordinal per variant) plus casts/binds.

### B. Mandatory dfVM (Implementation Details)

**Process model.** A dedicated JVM process per run; module graph stripped to the minimum; custom ClassLoader; strict classpath.

#### Capability enforcement.

- *Gates*: env, fs (virtual root/overlay, path policies), net (host/port ACL + HTTP fixture layer), time (virtual clock), rng (seeded PRNG).
- Mediation: stdlib calls (env/net/fs/time/rng/conc) go through dfVM services; Java interop calls are checked against the jvmAllowlist and capability guards (e.g., Http.get requires net).
- Deny: reflection, sun.misc.Unsafe, raw sockets, ProcessBuilder, arbitrary file IO outside virtual root.

#### Determinism & replay.

- Virtual *clock* and *rng* (seeded).
- HTTP fixtures: record bodies, headers, status; keyed by method+URL+headers+body; replay exact responses.

**Resource quotas.** CPU wallclock, heap soft cap, egress bytes, open FDs, concurrent tasks. Policy: fail-fast with trace.

**Tracing.** Every effect creates a span: name, start/stop, attributes (sizes, durations), contract verdicts, and optional previews (first N bytes of bodies).

**Packaging.** *dfpkg* (zip): bytecode jar, dfpkg.json manifest (entrypoint, capabilities, jvmAllowlist, hashes), optional fixtures, signature file. Loader verifies signature, hashes, and policy.

#### IV. CONCLUSION (PROPOSAL)

df++ aims to keep the surface small and semantics clear: interfaces over inheritance, functional collections over loops, explicit tasks with contracts, first-class pattern matching, and a mandatory sandbox. The ANTLR+ASM stack is straightforward; dfVM makes runs reproducible, auditable, and safe. The static solver catches simple issues early, and Java interop remains powerful while still capability-checked.

# APPENDIX A: MATHEMATICAL GRAMMAR (ABSTRACT SYNTAX)

We write a minimal abstract grammar (omitting obvious lists):

```
M
           ::= module p ; I^* ; D^*
Ι
           ::= \quad \mathsf{import} \ q \ \mathsf{as} \ x
           ::= \quad \mathsf{type} \ T \ = \ \Sigma \mid \mathsf{fn} \ f(\overline{x:T}): T \ = \ e \mid \mathsf{task} \ s \ \{C\} \mid \mathsf{const} \ x: T = e \mid \mathsf{let} \mid \mathsf{mut} \ x: T = e
D
\sum
           := K(\overline{T}) \mid \Sigma \mid \dots
           ::= c \mid x \mid e \text{ op } e \mid \lambda x : T. \ e \mid e(e) \mid \{\overline{x = e}\} \mid e.x \mid \operatorname{match}(e) \{\overline{\pi \to e}\}
e
           ::= K(\overline{x}) \mid \{\overline{x=y?}\} \mid c \mid \underline{\hspace{1cm}}
\pi
C
           ::= pre \phi; act \{S\} ; pos \psi
S
                     \overline{\text{stmt}}; e?
           ::=
                    propositional formulas with size, \leq, =, ...
\phi, \psi
```

Values and stores follow standard call-by-value with immutable const, single-assignment let, and mutable mut.

APPENDIX B: FORMAL SEMANTICS (TYPING & DYNAMICS)

**Typing.** We write  $\Gamma \vdash e : T$ .

$$\begin{split} \frac{\Gamma(x) = T}{\Gamma \vdash n : \text{Int}} & \frac{\Gamma(x) = T}{\Gamma \vdash x : T} & \frac{\Gamma \vdash e_1 : \text{Int} \quad \Gamma \vdash e_2 : \text{Int}}{\Gamma \vdash e_1 + e_2 : \text{Int}} \\ \text{Records:} & \frac{\forall i. \ \Gamma \vdash e_i : T_i}{\Gamma \vdash \{x_i = e_i\}_i : \{x_i : T_i\}_i}, & \text{field:} \\ \frac{\Gamma \vdash e : \{x : T, \ldots\}}{\Gamma \vdash e.x : T}. & \\ \text{Sum intro:} & \frac{\Gamma \vdash e : T}{\Gamma \vdash K(e) : K(T)}. & \text{Match:} \\ \frac{\Gamma \vdash e : \Sigma \quad \forall i. \ (\Gamma, \overline{x_i : T_i}) \vdash e_i : U \quad \text{exh}(\{\pi_i\}, \Sigma)}{\Gamma \vdash \text{match}(e)\{\pi_i \to e_i\}_i : U} \end{split}$$

where exh checks exhaustiveness (exact for finite  $\Sigma$ , warning otherwise).

**Tasks.** A task  $\tau$  evaluates as:

$$\frac{\Gamma \vdash \phi \Rightarrow \text{true}}{\tau.\text{pre} \Downarrow \text{true}} \qquad \frac{\tau.\text{pre} \Downarrow \text{true} \quad \{S\} \Downarrow v \quad \tau.\text{pos} \Downarrow \text{true}}{\tau \Downarrow \text{Ok}(v)}$$

Otherwise the result is Err(ContractFail). Orchestration is sequencing of these judgments; parallel evaluates in a product machine over Loom VTs and joins.

**Contracts.** Static solver reduces  $\phi, \psi$  to {Proved, Refuted, Unknown}; runtime always re-checks but may elide in *proved* mode.

# APPENDIX C: ANTLR GRAMMAR (UPDATED, V1 WITH 29 match) 30

31

```
lexer grammar DfppLexer;
2
    // keywords
    MODULE: 'module'; IMPORT: 'import'; AS: 'as';
   TYPE: 'type'; FN: 'fn'; TASK: 'task'; PRE: 'pre'; ACT: 'act'; POS: 'pos'; RUN: 'run'; PARALLEL: 'parallel';
    MATCH: 'match';
    CONST: 'const'; LET: 'let'; MUT: 'mut';
    TRUE:'true'; FALSE:'false'; NULL:'null';
10
11
    // literals / identifiers
12
    INTLIT:[0-9]+; DECLIT:[0-9]+'.'[0-9]+;
13
    STRING: '"' ( ~["\\] | '\\' . )* '"' | '\'' ( ~['\\] | '\\' . )* '\'' ;
14
15
    ID:[a-zA-Z_][a-zA-Z_0-9]*;
16
    BTICKID:'`' ( ~[`\\] | '\\' . )+ '`';
17
    // symbols
19
    LP:'('; RP:')'; LB:'{'; RB:'}'; LSB:'['; RSB:']'; COMMA:','; DOT:'.'; COLON:':'; EQ:'='; QMARK:'?';
20
21
          ARROW: '->';
    PLUS: '+'; MINUS: '-'; STAR: '*'; SLASH: '/'; PERC: '%';
22
    AND: '&&'; OR: '||'; NOT: '!'; EE: '=='; NE: '!='; LT: '<'; LE: '<='; GT: '>'; GE: '>=';
23
    // trivia
    WS:[ \t\n]+ -> skip;
26
    LINE_COMMENT:'//' ~[\r\n]* -> skip;
BLOCK_COMMENT:'/*' .*? '*/' -> skip;
```

```
parser grammar DfppParser;
   options { tokenVocab = DfppLexer; }
   program : moduleDecl? importDecl* topDecl* EOF ;
   moduleDecl: MODULE qid;
   importDecl: IMPORT (qid | JAVA qid) AS ID ; //
       proposal: 'import java: ... as ...'
           : ident (DOT ident)*;
   ident
             : ID | BTICKID ;
   topDecl : constDecl | letDecl | typeDecl | fnDecl |
       taskDecl /* | testDecl (tooling) */;
11
   constDecl : CONST ident (COLON typeRef)? EQ expr ;
12
   letDecl : (LET|MUT) ident (COLON typeRef)? EQ expr ;
13
   typeDecl : TYPE ident EQ sumType ;
15
           : variant ('|' variant)* ;
   sumType
           : ident (LP typeRef (COMMA typeRef)* RP)?;
17
   // tasks use STRING names to allow spaces
19
   taskDecl : TASK STRING LB (PRE expr)? ACT block (POS
20
       expr)? RB;
21
   // blocks and statements
22
            : LB stmt* expr? RB ;
23
            : constDecl | letDecl | runStmt |
   stmt
       parallelStmt | expr ;
   runStmt : runExpr :
   parallelStmt : PARALLEL LB stringList RB ;
   stringList: STRING (COMMA STRING)*;
```

```
// expressions
expr
 : expr OR expr
 I expr AND expr
 | expr EE expr | expr NE expr
 | expr LT expr | expr LE expr | expr GT expr | expr GE
 | expr PLUS expr | expr MINUS expr | expr STAR expr |
    expr SLASH expr | expr PERC expr
 | NOT expr
 | primary QMARK expr COLON expr
 | primary
primary
 : literal
 | ident
 | call
 I memberChain
 | LP expr RP
 | recordLit
 | arrayLit
 | matchExpr
         : INTLIT | DECLIT | STRING | TRUE | FALSE |
literal
    NULL ;
          : ident LP args? RP ;
call
          : expr (COMMA expr)*;
args
// run('A').then('B') as a chain
runExpr
        : memberChain ;
memberChain
: runCall (DOT ident LP args? RP )*
runCall : RUN LP STRING RP ;
// pattern matching
matchExpr : MATCH LP expr RP LB matchArm+ RB ;
matchArm : pattern ARROW expr ;
pattern : ident (LP (ident (COMMA ident)*)? RP)?
    K(x,y) or ctor()
         | LB patField (COMMA patField)* RB
                                                      //
    \{k = x, \ldots\}
          | literal
patField : ident EQ ident ;
recordLit : LB recField (COMMA recField)* RB ;
recField : ident EQ expr ;
arrayLit : LSB (expr (COMMA expr)*)? RSB ;
         : qid | LB typeField (COMMA typeField)* RB |
    typeRef ARROW typeRef ;
typeField : ident COLON typeRef ;
```

# APPENDIX D: STANDARD LIBRARY INTERFACES (V1, FULL Types)

Implementations live behind dfVM; these are the signatures. **std.core** 

- type Option<T> = None | Some(T)
- type Result<T,E> = Ok(T) | Err(E)
- trait Future<T> { map<U>((T)->U): Future<U>; then<U>((T)->Future<U>): Future<U>; await(): T }
- List<T>:

  - fold<A>(A, (A,T)->A): A, reduce((T,T)->T): T,
     scan<A>(A,(A,T)->A): List<A>
  - zip<U>(List<U>): List<(T,U)>, take(Int): List<T>, drop(Int): List<T>
  - groupBy<K>((T)->K): Map<K,List<T»,
     sortBy<K:Comparable>((T)->K): List<T>
  - forEach((T)->Unit): Unit, size(): Int, isEmpty():
    Bool, nonEmpty(): Bool
  - static range(Int, Int): List<Int>
- Set<T>:
  - union(Set<T>): Set<T>, intersect(Set<T>):
     Set<T>, diff(Set<T>): Set<T>
  - contains(T): Bool, size(): Int, isEmpty(): Bool
  - map<U>((T)->U): Set<U>, filter((T)->Bool):
     Set<T>, toList(): List<T>
- Map<K,V>:
  - get(K): Option<V>, put(K,V): Map<K,V>, remove(K):
     Map<K,V>
  - containsKey(K): Bool, keys(): Set<K>, values():
     List<V>, entries(): List<(K,V)>

  - size(): Int, isEmpty(): Bool
- Option.match<U>(some:(T)->U, none:()->U): U
- Result.match<U>(ok:(T)->U, err:(E)->U): U
- getOrElse(T): T on Option and Result<T,E>
- Base: all base types expose str(): String

#### std.json

- type Json = ... (opaque)
- parse(String): Result<Json,String
- stringify(Json): String
- encode<T>(T): Json (derived for records/ADTs)
- decode<T>(Json): Result<T,String

std.io — print(String): Unit

# std.sys.env/time/rng/net/conc (via dfVM)

- env: env(String): Option<String>
- time: nowMs(): Long, sleep(Long): Unit
- rng: randBytes(Int): Bytes
- net:
  - type HttpResp = { status: Int, headers Map<String>, body: Bytes }
  - httpGet(String, Map<String,String>): HttpResp

- httpPost(String, Map<String,String>, Bytes):
   HttpResp
- conc: spawn(()->Unit): Future<Unit>

#### std.flow.orch

- type TaskChain = { then(String): TaskChain, await(): Result<Unit,String> }
- run(String): TaskChain
- parallel(List<String>): Result<Unit,String>

std.contracts — assert(Bool): Unit; pipelines accept
.inv(Bool) (debug-mode check).

std.jvm (only if the direct Java import shortcut is not used)

- callStatic(class:String, method:String, paramSig:String, args:List<Any>): Any
- new(class:String, paramSig:String, args:List<Any>): Any

#### APPENDIX E: WORKED EXAMPLES

```
Listing 1. ::df++ — Pattern matching with lists and options
   ::df++
2
   module demo.match
   type Option<T> = None | Some(T)
   fn head<T>(xs: List<T>): Option<T> =
     match (xs.size()) {
       0 -> None
        _ -> Some(xs.take(1).head)
                                               // assume head
10
        on non-empty
11
12
   fn greet(opt: Option<String>): String =
13
     match (opt) {
  Some(s) -> "Hello, " + s
14
15
                -> "Hello"
       None
16
```

```
Listing 2. ::df++ — Tasks, sequencing, parallelism
   ::df++
2
   module demo.tasks
   import std.json as Json
   import std.sys.net as Net
   import std.io as IO
   type Post = { id: Int, title: String }
   mut posts: List<Post> = []
10
   mut titles: List<String> = []
11
12
   task 'fetch posts' {
13
     act {
14
       let r = Net.httpGet("https://example.org/posts", {})
       posts = Json.parse(r.body.utf8())
16
                   .then(j -> Json.decode<List<Post>>(j))
17
                   .getOrElse([])
18
19
     pos posts.size() >= 0
20
21
22
   task 'process titles' {
23
     pre posts.nonEmpty()
     act {
25
       titles = posts.filter(p ->
        p.title.toLower().contains("ai"))
                       .map(p -> p.title)
27
28
29
     pos titles.size() <= posts.size()</pre>
30
31
   fn main(): Unit = {
32
     run('fetch posts').then('process titles')
parallel { 'fetch posts', 'process titles' }
33
     IO.print("titles: " + titles.size().str())
35
```

# APPENDIX F: TOOLING (PROPOSAL: CLI, LSP, TEST, PACKAGING)

**CLI.** dfpp build  $\rightarrow$  signed .dfpkg; dfpp run; dfpp test (discovers test '...' blocks); dfpp trace; dfpp replay.

**LSP.** Hover types, go-to-def, rename, contract diagnostics (static solver verdict), quick-fixes for missing types/capabilities.

**Debugger.** Step tasks, inspect contract verdicts, see dfVM caps, tail spans.

**dfpkg format.** Zip with: bytecode JAR, dfpkg.json (entrypoint, capabilities, jvmAllowlist, hashes), fixtures/, bundle.sig. Loader verifies signature + policy.

**Tests.** test '...' blocks run under dfVM with deterministic time/RNG; JUnit XML output for CI.

# APPENDIX G: CASE STUDY — FLOWTOMIC DEPLOYMENT (EXPANDED)

#### Mapping nodes to df++.

- Each node → task 'node name'; validation → pre; success criteria → pos.
- Edges → run('A').then('B') or parallel{ 'A', 'B'}.
- Shared transforms  $\rightarrow$  pure fns; data reshaping uses List/Set/Map.

**External code** (**TS/Python**). For v1, Flowtomic "code actions" run in sidecar services (Node/Python). df++ tasks call them via std.sys.net.http\* (record/replay supported), keeping dfVM deterministic and safe. (A GraalVM plugin is future work.)

# dfpkg manifest (example).

```
"entry": "example.flow.Main.main",
2
3
      "capabilities": {
        "net": { "hosts": ["api.example.org:443"] },
"env": ["API_KEY"],
        "time": { "deterministic": true },
        "rng": { "seed": 42 },
"fs": { "root": "/work", "writable": ["/work/out"] }
     },
"jvmAllowlist": [
10
11
        com.vendor.sdk.Client#get(Ljava/lang/String;)Ljava/lang/String;"
12
     "fixtures": "fixtures/",
13
      "hashes": { "jar": "sha256:...", "fixtures": "
14
        sha256:..."},
      "signature": "ed25519:..."
15
   }
16
```

# Operational flow.

- 1) Author designs graph, exports to df++.
- 2) dfpp build produces a signed .dfpkg.
- 3) Runtime loads under dfVM, which enforces capabilities, traces effects, and optionally records fixtures.
- 4) CI uses dfpp replay to reproduce runs bit-for-bit (virtual time/RNG, HTTP fixtures).

Why this pairing works. Flowtomic's graph abstractions align directly with tasks+orchestration; data wrangling uses functional collections; dfVM provides operational guarantees (least-privilege, determinism, audit) crucial for deployed models.